



**BRENT OZAR**  
UNLIMITED®

# Fundamentals of Index Tuning

## Part 3: Indexing for ORDER BY

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)

O3 p1

## Agenda

How ORDER BY comes into play

Combining WHERE and ORDER BY

TOP exceptions: when ORDER BY goes first

How parameters affect key order



**I'd like to place an  
ORDER BY  
after two equality searches**

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)



03 p3

## Bring some order around here

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;
```



## Think back to your 2 earlier indexes.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;

CREATE INDEX IX_DisplayName_Location
ON dbo.Users(DisplayName, Location);

CREATE INDEX IX_Location_DisplayName
ON dbo.Users(Location, DisplayName);
```



## Add new versions with Reputation

```
CREATE INDEX IX_DisplayName_Location_Reputation
ON dbo.Users(DisplayName, Location, Reputation);

CREATE INDEX IX_Location_DisplayName_Reputation
ON dbo.Users(Location, DisplayName, Reputation);

/* Plus a third idea: */
CREATE INDEX IX_Reputation_DisplayName_Location
ON dbo.Users(Reputation, DisplayName, Location);
```



```
SET STATISTICS IO ON;
GO
SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = 1) /* Clustered index scan */
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Reputation)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_Location_DisplayName_Reputation)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_Reputation_DisplayName_Location)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;
GO
```

**Test 'em**

O3 p7

## Survey says...

| Index                              | Logical Reads | Total Pages in the Index |
|------------------------------------|---------------|--------------------------|
| Clustered index (white pages)      | 45,184        | 45,184                   |
| IX_DisplayName_Location_Reputation | 4             | 13,995                   |
| IX_Location_DisplayName_Reputation | 4             | 14,486                   |
| IX_Reputation_DisplayName_Location | 13,996        | 13,996                   |

Ouch. Putting reputation first meant no seeking at all, and we scanned the whole thing. (Still better than a table scan though.)






## Which one does SQL Server pick?

```
90  /* Which one does SQL Server pick? */
91  SELECT Id, DisplayName, Location
92  FROM dbo.Users
93  WHERE DisplayName = 'alex'
94  AND Location = 'Seattle, WA'
95  ORDER BY Reputation;
96  GO
```

The one that leads with  
DisplayName.



# ORDER BY

after an INequality search

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)



03 p10

## Your last query:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;
```



## Let's go anywhere BUT Seattle

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

What's the perfect index for this?  
How selective is each part of the filter?



## Survey says...

| Index                              | Logical Reads | Total Pages in the Index |
|------------------------------------|---------------|--------------------------|
| Clustered index (white pages)      | 45,184        | 45,184                   |
| IX_DisplayName_Location_Reputation | 13            | 13,995                   |
| IX_Location_DisplayName_Reputation | 4,864         | 14,486                   |
| IX_Reputation_DisplayName_Location | 13,996        | 13,996                   |

Ouch. Putting reputation first meant no seeking at all, and we scanned the whole thing. (Still better than a table scan though.)



## So the perfect index for it:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

```
CREATE INDEX IX_DisplayName_Location_Reputation ON
dbo.Users (DisplayName, Location, Reputation);
```

Step 1: seek to Alex

Step 2: scan through,  
returning everyone  
EXCEPT Seattle

Step 3: read them out  
sorted by Reputation,  
except...they're not.


Logistics, chat, questions, recording  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)

O3 p14

## Our index gets used, but...

```
110 SELECT Id, DisplayName, Location
111     FROM dbo.Users
112     WHERE DisplayName = 'alex'
113           AND Location <> 'Seattle, WA'
114     ORDER BY Reputation;
115 GO
```

The plan has a Sort even though the data in the index is sorted in order – isn't it?



## Write a query to visualize the index

```
CREATE INDEX IX_DisplayName_Location_Reputation ON  
dbo.Users (DisplayName, Location, Reputation);
```

```
SELECT DisplayName, Location, Reputation, Id  
FROM dbo.Users  
ORDER BY DisplayName, Location, Reputation;
```

150 %

| Results |             |                        |            |        | Messages | Execution plan |
|---------|-------------|------------------------|------------|--------|----------|----------------|
|         | DisplayName | Location               | Reputation | Id     |          |                |
| 1       | quico       | London, United Kingdom | 14747      | 389099 |          |                |
| 2       | µBio        | California             | 8999       | 9796   |          |                |
| 3       | µlrad       | Tehran, Iran           | 5          | 136691 |          |                |
| 4       | 0_          | NULL                   | 48302      | 515054 |          |                |
| 5       | 0_o         | NULL                   | 1          | 418884 |          |                |
| 6       | 0_o         | NULL                   | 1          | 438437 |          |                |
| 7       | 0_o         | NULL                   | 3          | 406169 |          |                |

Logistics  
BrentOzar

03 p16



## Seek down to Alex

```
SELECT DisplayName, Location, Reputation, Id
FROM dbo.Users
ORDER BY DisplayName, Location, Reputation;
```

|     | DisplayName | Location         | Reputation | Id     |
|-----|-------------|------------------|------------|--------|
| 3.. | Arieltibby  | Japan            | 2191       | 405432 |
| 3.. | alotao      | Paris, France    | 2099       | 289493 |
| 3.. | xixiang     | Guangzhou, China | 4977       | 54140  |
| 3.. | alvando     | NULL             | 16         | 169640 |
| 3.. | Alex        | NULL             | 1          | 72886  |
| 3.. | Alex        | NULL             | 1          | 190109 |
| 3.. | Alex        | NULL             | 1          | 193346 |
| 3.. | Alex        | NULL             | 1          | 202170 |
| 3.. | Alex        | NULL             | 1          | 210507 |
| 3.. | Alex        | NULL             | 1          | 213524 |
| 3.. | Alex        | NULL             | 1          | 213688 |
| 3.. | Alex        | NULL             | 1          | 216795 |
| 3.. | Alex        | NULL             | 1          | 218333 |
| 3.. | Alex        | NULL             | 1          | 220772 |
| 3.. | Alex        | NULL             | 1          | 232257 |
| 3.. | Alex        | NULL             | 1          | 233160 |
| 3.. | Alex        | NULL             | 1          | 236336 |
| 3.. | Alex        | NULL             | 1          | 237014 |

Remember, we need them ordered by Reputation.

At first it looks like this will work, but...



```
SELECT DisplayName, Location, Reputation, Id
FROM dbo.Users
ORDER BY DisplayName, Location, Reputation;
```

| 10 %        |                           |                |        |
|-------------|---------------------------|----------------|--------|
| Results     | Messages                  | Execution plan |        |
| DisplayName | Location                  | Reputation     | Id     |
| Alex        |                           | 201            | 215550 |
| Alex        |                           | 200            | 200005 |
| Alex        |                           | 1133           | 94019  |
| Alex        |                           | 1189           | 91631  |
| Alex        |                           | 1818           | 79962  |
| Alex        |                           | 13733          | 83185  |
| Alex        | "The Cloud"               | 405            | 231736 |
| Alex        | Abenachod, India          | 161            | 419711 |
| Alex        | Ableside, NC              | 3              | 83742  |
| Alex        | Amsterdam, The Netherl... | 121            | 464950 |
| Alex        | Amway, France             | 96             | 383173 |
| Alex        | Athens, Greece            | 563            | 423581 |
| Alex        | Atlanta, GA               | 1493           | 237090 |
| Alex        | Auckland, New Zealand     | 6              | 414011 |
| Alex        | Austin, TX                | 140            | 435381 |
| Alex        | Austin, TX                | 240            | 177746 |
| Alex        | Australia                 | 817            | 167011 |
| Alex        | Australia                 | 1473           | 211869 |
| Alex        | Austria                   | 1935           | 276837 |
| Alex        | Barcelona, Spain          | 4967           | 25787  |
| Alex        | Bay Area, CA, United S... | 103            | 389684 |
| Alex        | Belarus                   | 137            | 514389 |
| Alex        | Belarus                   | 583            | 19081  |
| Alex        | Belgium                   | 81             | 422522 |
| Alex        | Bermuda                   | 1              | 363997 |
| Alex        | Billinghurst, United K... | 1364           | 181739 |
| Alex        | Boca Raton, FL            | 5725           | 257404 |
| Alex        | Borde, Sweden             | 71             | 454570 |
| Alex        | Boston, MA                | 1130           | 496645 |
| Alex        | Boston, MA                | 1965           | 417291 |
| Alex        | Brazil, Brazil            | 1825           | 312808 |
| Alex        | Brisbane, The Netherlands | 329            | 236813 |
| Alex        | Brighton, United Kingd... | 53             | 289608 |
| Alex        | Brooklyn, NY              | 1231           | 77204  |

## Reputation isn't sorted.

We're going to skip everyone who isn't in Seattle.

That means we need all the Alexes on this screen, plus more.

And they're not sorted by Reputation.

The fact that Reputation is "sorted" isn't helping here.



O3 p18

## Ordering Reputation doesn't help.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

```
CREATE INDEX IX_DisplayName_Location_Reputation ON
dbo.Users (DisplayName, Location, Reputation);
```

Step 1: seek to Alex

Step 2: scan through,  
returning everyone  
EXCEPT Seattle

Step 3: read them out  
sorted by Reputation,  
except...they're not.

Logistics, chat, questions, recording  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)

03 p19

## To prove it, create another index:

```
CREATE INDEX IX_DisplayName_Location_Reputation  
ON dbo.Users  
(DisplayName, Location, Reputation);  
  
CREATE INDEX IX_DisplayName_Location_Includes ON  
dbo.Users  
(DisplayName, Location) INCLUDE (Reputation);
```



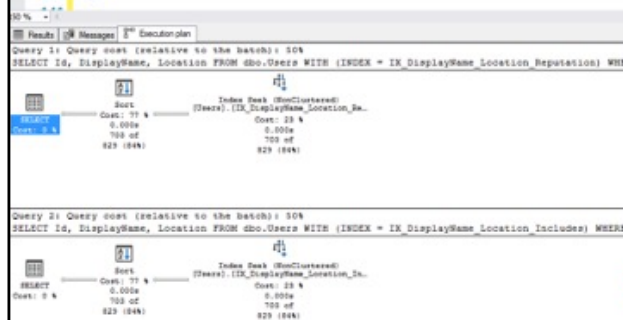
## They both get the same plan

```
129 SELECT Id, DisplayName, Location
130 FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Reputation)
131 WHERE DisplayName = 'alex'
132 AND Location <> 'Seattle, WA'
133 ORDER BY Reputation;
134
135 SELECT Id, DisplayName, Location
136 FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Includes)
137 WHERE DisplayName = 'alex'
138 AND Location <> 'Seattle, WA'
139 ORDER BY Reputation;
140 GO
```

Use an index hint to test both indexes separately.

Both do the sort.

And both have the same number of logical reads.



## Inequality searches make it tricky.

```
WHERE DisplayName = 'alex'  
      AND Location <> 'Seattle, WA'  
ORDER BY Reputation;
```

After you do an inequality search on a field, the sorting of subsequent fields in the index are usually less useful.

(That's a mouthful.)



## Putting Reputation SECOND helps.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;

CREATE INDEX IX_DisplayName_Reputation_Location ON
dbo.Users (DisplayName, Reputation, Location);
```

Step 1: seek to Alex

Step 2: the sort isn't  
needed: they're sorted

Step 3: Skip the users  
who aren't in Seattle

Logistics, chat, questions, recording  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)

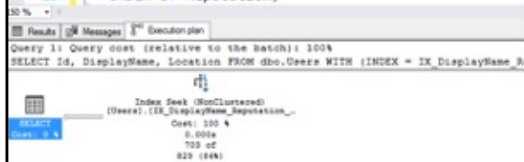
03 p23

## The sort is gone with this trick.

```
144 /* Promote Reputation one level: */
145 CREATE INDEX IX_DisplayName_Reputation_Location
146 ON dbo.Users(DisplayName, Reputation, Location);
147 GO
148
149 /* And the sort is gone: */
150 SELECT Id, DisplayName, Location
151 FROM dbo.Users WITH (INDEX = IX_DisplayName_Rep
152 WHERE DisplayName = 'alex'
153 AND Location <> 'Seattle, WA'
154 ORDER BY Reputation;
```

Obscure trick. To get it, key on:

1. Equality fields, then
2. Sort fields, then
3. Inequality fields



Query 1: Query cost (relative to the batch): 100%

SELECT Id, DisplayName, Location FROM dbo.Users WITH (INDEX = IX\_DisplayName\_Reputation\_Location) WHERE DisplayName = 'alex' AND Location <> 'Seattle, WA' ORDER BY Reputation;

| Index Scan (NonClustered)                   |
|---|
| (Users): IX_DisplayName_Reputation_Location |
| Cost: 100 %                                 |
| 0.000s                                      |
| 703 cr                                      |
| 829 (84%)                                   |






## SQL Server picks it, too.

```
157 /* Which one does SQL Server pick? */
158 SELECT Id, DisplayName, Location
159 FROM dbo.Users
160 WHERE DisplayName = 'alex'
161 AND Location <> 'Seattle, WA'
162 ORDER BY Reputation;
163 GO
164
```

If we don't hint the query, here it picks the index that removes the sort.



## What we've learned so far

Indexes help by pre-sorting rows to prep them for:

- WHERE: finding the rows we want
- ORDER BY: sorting them on the way out the door
- GROUP BY, FROM, JOINS, CTEs:  
more on these later

And so far, it kinda seems like you want to put keys in that same order: WHERE first, then ORDER BY. But that's not exactly how it works.



**S**  
**TOP**

**me if you've heard this one before**

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)



**O3 p27**

## Start by dropping your indexes.

We're going to tackle a new set of queries, and I don't want to confuse SQL Server's hints with existing indexes.

```
EXEC DropIndexes;
```

Get the code:

[BrentOzar.com/go/dropindexes](https://BrentOzar.com/go/dropindexes)



## Design an index for this:

```
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > 1  
ORDER BY CreationDate ASC;
```



## Which field should we lead with?

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate
ON dbo.Users(Reputation, CreationDate);
```


```
CREATE INDEX IX_CreationDate_Reputation
ON dbo.Users(CreationDate, Reputation);
```



## If we lead with Reputation...

We seek to 2, but then we find 1.4M users that match!  
We have to sort 'em all by CreationDate.

```
187 SELECT TOP 100 Id, Reputation, CreationDate
188 FROM dbo.Users WITH (INDEX = IX_Reputation_CreationDate)
189 WHERE Reputation > 1
190 ORDER BY CreationDate ASC;
```



Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)



O3 p31

## Visualize the index contents

```
/* Visualizing the Reputation, CreationDate index: */  
SELECT Reputation, CreationDate, Id  
FROM dbo.Users  
ORDER BY Reputation, CreationDate;
```

|       | Reputation | CreationDate            | Id     |
|-------|------------|-------------------------|--------|
| 40711 | 1          | 2010-12-31 23:42:32.577 | 559546 |
| 40712 | 1          | 2010-12-31 23:46:14.137 | 559547 |
| 40713 | 1          | 2010-12-31 23:55:40.489 | 559550 |
| 40714 | 2          | 2008-09-15 18:17:21.780 | 5284   |
| 40715 | 2          | 2008-11-24 20:50:49.057 | 40501  |
| 40716 | 2          | 2008-11-26 09:59:50.243 | 40566  |
| 40717 | 2          | 2009-02-03 02:26:11.907 | 61757  |
| 40718 | 2          | 2009-02-25 13:51:56.000 | 70656  |
| 40719 | 2          | 2009-03-02 09:52:37.567 | 72731  |
| 40720 | 2          | 2009-06-24 11:01:29.467 | 126149 |
| 40721 | 2          | 2009-09-20 08:21:00.400 | 164710 |
| 40722 | 2          | 2009-09-20 10:13:40.479 | 164767 |
| 40723 | 2          | 2009-10-10 10:12:12.143 | 189994 |
| 40724 | 2          | 2009-10-14 20:04:56.193 | 190136 |

When the index is on Reputation, CreationDate, we can seek to 2, but...are the first 10 users we find the lowest CreationDates overall?

Or just the lowest for Reputation = 2?





## Easier way to see it

```
/* Visualizing the Reputation, CreationDate, Id  
SELECT Reputation, CreationDate, Id  
FROM dbo.Users  
ORDER BY Reputation, CreationDate;
```

|        | Reputation | CreationDate            | Id     |
|--------|------------|-------------------------|--------|
| 277696 | 3819       | 2010-03-25 08:32:56.610 | 301514 |
| 277697 | 3819       | 2010-06-14 02:20:25.290 | 365977 |
| 277698 | 3819       | 2010-07-19 15:48:21.980 | 395875 |
| 277699 | 3820       | 2009-09-24 08:03:37.333 | 21537  |
| 277700 | 3820       | 2009-07-30 09:26:28.910 | 147695 |
| 277701 | 3820       | 2010-10-08 09:07:30.453 | 470062 |
| 277702 | 3821       | 2009-10-05 08:00:20.540 | 25234  |
| 277703 | 3821       | 2008-12-30 00:22:36.087 | 50025  |
| 277704 | 3821       | 2010-04-12 15:33:11.823 | 314670 |
| 277705 | 3821       | 2010-12-29 08:26:55.083 | 556899 |
| 277706 | 3822       | 2009-03-25 01:36:44.977 | 82333  |
| 277707 | 3822       | 2010-01-22 20:12:36.790 | 257065 |
| 277708 | 3822       | 2010-09-01 10:51:24.297 | 436853 |

It's more obvious when we page down to higher Reputation numbers.

The CreationDate keeps resetting with each new Reputation.

The sort on the second field is less useful when we're scanning.



## What if we lead with CreationDate?

```
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > 1  
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate  
ON dbo.Users(Reputation, CreationDate);
```



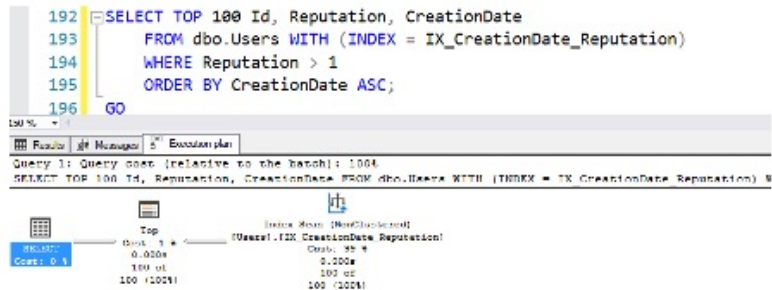
```
CREATE INDEX IX_CreationDate_Reputation  
ON dbo.Users(CreationDate, Reputation);
```



## We “scan” the index, but...

Remember from How to Think Like the Engine: scan just means we start at one end of the index, and we read until we find the rows that match.

And there's no sort! The data is already sorted.



## Visualize the index contents

```
/* Visualizing the Reputation, CreationDate index: */  
SELECT CreationDate, Reputation, Id  
FROM dbo.Users  
ORDER BY CreationDate, Reputation
```

| Results                 | Msg        | Msgpack | W | Execution plan |
|-------------------------|------------|---------|---|----------------|
| CreationDate            | Reputation | Id      |   |                |
| 2008-07-31 00:00:00.000 | 1          | 1       |   |                |
| 2008-07-31 14:22:31.287 | 3491       | 2       |   |                |
| 2008-07-31 14:22:31.287 | 13418      | 3       |   |                |
| 2008-07-31 14:22:31.287 | 44300      | 4       |   |                |
| 2008-07-31 14:22:31.317 | 28768      | 4       |   |                |
| 2008-07-31 14:22:31.317 | 39172      | 5       |   |                |
| 2008-07-31 21:00:24.057 | 942        | 8       |   |                |
| 2008-07-31 21:08:06.517 | 14337      | 9       |   |                |
| 2008-07-31 21:57:06.240 | 101        | 10      |   |                |
| 2008-08-01 00:05:11.147 | 1890       | 11      |   |                |
| 2008-08-01 04:18:04.843 | 177138     | 12      |   |                |
| 2008-08-01 12:01:03.661 | 41167      | 13      |   |                |
| 2008-08-01 12:01:03.023 | 527        | 16      |   |                |
| 2008-08-01 12:02:21.617 | 44443      | 17      |   |                |
| 2008-08-01 12:05:14.233 | 1272       | 18      |   |                |
| 2008-08-01 12:06:11.070 | 1600       | 20      |   |                |
| 2008-08-01 12:11:11.887 | 12816      | 22      |   |                |
| 2008-08-01 12:11:43.703 | 4296       | 23      |   |                |
| 2008-08-01 12:12:03.473 | 3031       | 24      |   |                |
| 2008-08-01 12:18:23.243 | 16381      | 25      |   |                |
| 2008-08-01 12:18:14.640 | 11411      | 26      |   |                |
| 2008-08-01 12:21:40.020 | 1231       | 27      |   |                |
| 2008-08-01 12:24:53.820 | 72928      | 29      |   |                |

When the index is on  
CreationDate, Reputation, we  
start reading, looking for 100  
users with Reputation > 1.

They almost all match!

As soon as we read 100 rows  
that match, we're done. No  
need to scan the whole index.



## Survey says...

| Index                         | Logical Reads | Total Pages in the Index |
|-------------------------------|---------------|--------------------------|
| Clustered index (white pages) | 45,184        | 45,184                   |
| IX_Reputation_CreationDate    | 3,805         | 6,812                    |
| IX_CreationDate_Reputation    | 3             | 6,817                    |



**In this case, the ORDER BY field should go first in the index.**

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate
ON dbo.Users(Reputation, CreationDate);
```

```
CREATE INDEX IX_CreationDate_Reputation
ON dbo.Users(CreationDate, Reputation);
```





**Remember  
selectivity?**



## TOP is kinda like a WHERE clause.

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

That's kinda like saying:

```
SELECT stuff
FROM dbo.Users
WHERE (user is in the top ~100) by CreationDate
```





## So let's keep just this one for now

```
DropIndexes;
```

```
CREATE INDEX IX_CreationDate_Reputation  
ON dbo.Users(CreationDate, Reputation);
```

Let's say we decided to just keep this one.



## Now run this.

```
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > 1000000  
ORDER BY CreationDate ASC;
```

There aren't a lot of rows with  
Reputation > 1,000,000.



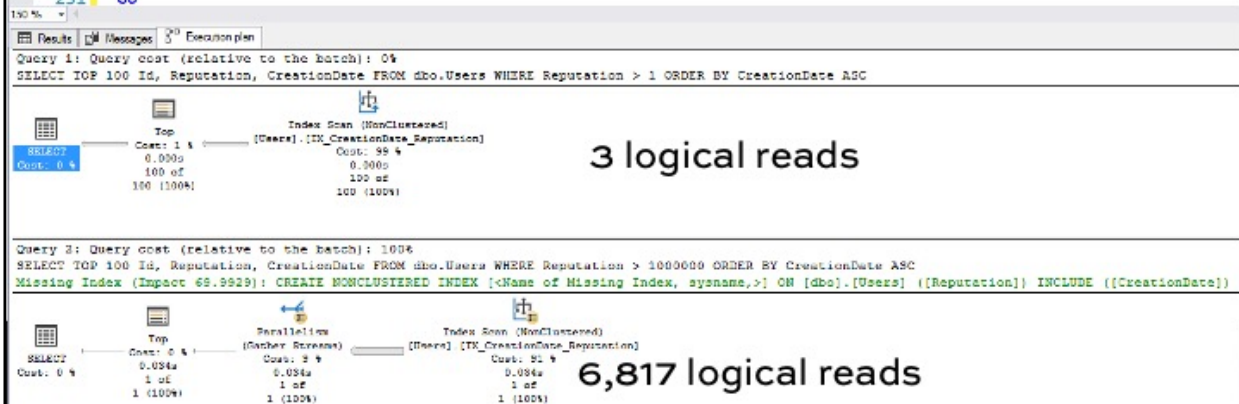
```

220  /* The original query: */
221  SELECT TOP 100 Id, Reputation, CreationDate
222      FROM dbo.Users
223      WHERE Reputation > 1
224      ORDER BY CreationDate ASC;
225
226  /* The new one looking for Jon Skeet: */
227  SELECT TOP 100 Id, Reputation, CreationDate
228      FROM dbo.Users
229      WHERE Reputation > 1000000
230      ORDER BY CreationDate ASC;
231  GO

```

**Both old & new queries  
use the index...**

But the index isn't as good of a  
fit for the second query. Why?



## Jon Skeet isn't in the first 100.

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1000000
ORDER BY CreationDate ASC;
```


The TOP 100 by CreationDate is only selective IF the person you're looking for is in that list.

In this case, WHERE Reputation > 1000000 is much more selective – that should go first.









Index for the WHERE  
to reduce reads

Index for the ORDER BY  
to reduce sorts

**Indexing requires  
compromises and choices.**

## These are just 2 inequality searches.

```
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > 1000000  
ORDER BY CreationDate ASC;
```

It comes down to:

- Which ones are the most selective
- And whether you want to cut reads or cut sorts
- Which parameters run the most often



## Say this is a stored procedure.

```
CREATE PROC usp_SearchUsers  
    @SearchReputation INT AS  
  
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > @SearchReputation  
ORDER BY CreationDate ASC;  
GO
```





```
CREATE PROC usp_SearchUsers  
    @SearchReputation INT AS  
  
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > @SearchReputation  
ORDER BY CreationDate ASC;  
GO
```

When @SearchReputation = 1, lots of data matches,  
so it's better to index on CreationDate, then Reputation.

When @SearchReputation = 1,000,000, then only 1 person matches,  
so it's better to index on Reputation, then CreationDate.





**Re-cap**

## Recap

If your WHERE clause is filtering just for equalities, then add the ORDER BY fields into the index key, and the index will handle all the sorting for you.

Out here in the real world, though, your query will have a mix of equality and inequalities.

Different parameter values affect key order too.

Our goal: get a good enough combination of keys to cover as many queries as practical.





**BRENT OZAR**  
UNLIMITED®

# Fundamentals of Index Tuning

Part 4: let's see what you learned about ORDER BY.

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)

03 p52

## Lab requirements

Download any Stack Overflow database:

- [BrentOzar.com/go/querystack](http://BrentOzar.com/go/querystack)
- I'm using the 50GB Stack Overflow 2013 (but any year is fine, even the 10GB one)

Desktop/laptop requirements:

- Any supported SQL Server version will work
- The faster your machine, the faster your indexes will get created



## Working through the lab

Read the first query, execute it, do your work inline, taking notes as you go

90 minutes: you work through the lab, asking questions in Slack as you go, and get lunch (either lunch first, or after your work)

The live stream will be off during lunch.

After lunch: I work through it onscreen

Logistics, chat, questions, recording info:  
[BrentOzar.com/training/live](https://BrentOzar.com/training/live)



03 p54