



BRENT OZAR
UNLIMITED®

How Indexes Help Avoid Blocking

2.4 p1

Concurrency challenges

Locking: Lefty takes out a lock.

Blocking: Righty needs a lock, but Lefty has it.
SQL Server will let Righty wait for forever,
and the symptom is LCK* waits.

Deadlocks:

Lefty has locks, but needs some held by Righty.
Righty has locks, but needs some held by Lefty.
SQL Server solves this one by killing somebody,
and the symptom is dead bodies everywhere.



2.4 p2

3 ways to fix blocking & deadlocks

1. Have enough indexes to make your queries fast, but not so many that they slow down DUIs, making them hold more locks for longer times.
(This session focuses on this.)
2. Keep batches & transactions short and sweet.
(We cover this in Mastering Query Tuning.)
3. Use the right isolation level for your app's needs.
(We cover this in Mastering Server Tuning.)

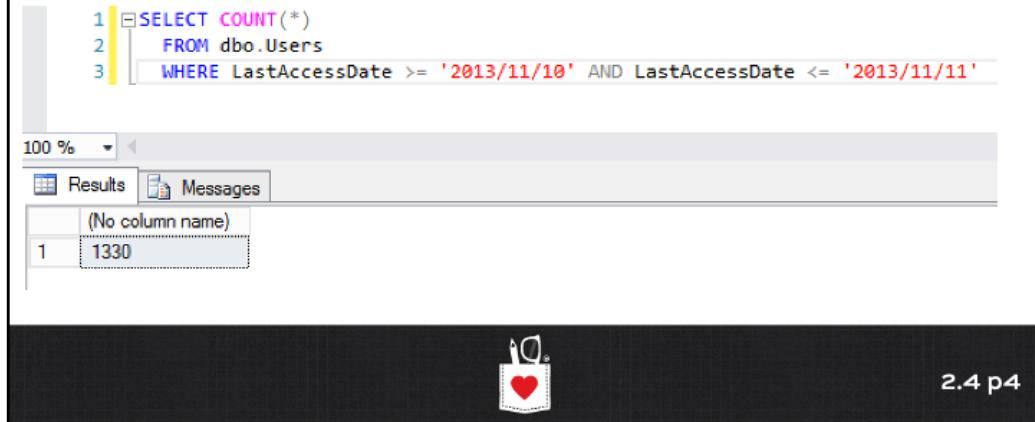


2.4 p3

Let's go back to the Users table.

You only have the clustered index on ID,
the white pages of the table.

How many accessed the system on my birthday?



The screenshot shows a SQL Server Management Studio (SSMS) interface. In the query editor, a T-SQL script is displayed:

```
1 | SELECT COUNT(*)
2 |   FROM dbo.Users
3 | WHERE LastAccessDate >= '2013/11/10' AND LastAccessDate <= '2013/11/11'
```

The results pane shows a single row of data:

| | (No column name) |
|---|------------------|
| 1 | 1330 |

Below the results pane, there is a small decorative icon of a red heart inside a white box with a black border, and the text "2.4 p4".

Let's reward them.

You only have the clustered index on ID,
the white pages of the table.

What's the execution plan for this query:

```
BEGIN TRAN
UPDATE dbo.Users
SET Reputation = Reputation + 100
WHERE LastAccessDate >= '2013/11/10'
AND LastAccessDate <= '2013/11/11'
```

dbo.Users - Clustered Index

| Id | Rep | CreationDate | DisplayName | LastAccessDate | Location | Age | AboutMe |
|-----|------|------------------|--------------|------------------|----------------|------|------------------------------|
| 1 | 2406 | 7/12/09 10:51 PM | Jeff Atwood | 4/1/10 10:35 AM | El Cerrito, CA | 39 | = '1800/01/01'
AND LastAccessDate <= '1800/01/02'
```

dbo.Users - Clustered Index

| Id  | Rep  | CreationDate     | DisplayName  | LastAccessDate   | Location       | Age  | AboutMe                     |  |  |
|-----|------|------------------|--------------|------------------|----------------|------|-----------------------------|--|--|
| 1   | 2406 | 7/12/09 10:51 PM | Jeff Atwood  | 4/1/10 10:35 AM  | El Cerrito, CA | 39   | = '1800/01/01'  
4 | AND LastAccessDate <= '1800/01/02'  
5 |  
  
100 % < Messages Execution plan  
Query 1: Query cost (relative to the batch): 100%  
SELECT Id FROM dbo.Users WHERE LastAccessDate >= '1800/01/01' AND La  
Missing Index (Impact 94.2875): CREATE NONCLUSTERED INDEX [<Name of  
[Icon] [Icon] [Icon]  
SELECT Cost: 0 % Parallelism (Gather Streams) Clustered Index Scan (Clustered)  
[Users].[PK_Users_Id] Cost: 4 % Cost: 96 %
```

We don't have an index on LastAccessDate,
so we're going to have to scan the table for this.

If we run it, what happens?



2.4 p9

Our SELECT just sits there blocked.

```
SQLQuery1.sql - X
1 BEGIN TRAN
2 UPDATE dbo.Users
3 SET Reputation = Reputation + 100
4 WHERE LastAccessDate >= '2013/11/10'
5 AND LastAccessDate <= '2013/11/11'

100 % < Messages
(1330 row(s) affected)

SQLQuery2.sql - Executing... - X
1 SELECT Id
2 FROM dbo.Users
3 WHERE LastAccessDate >= '1800/01/01'
4 AND LastAccessDate <= '1800/01/02'
5

100 % < Results Messages
```

The update on the left has locks on some rows in the Users table.

Until that lock is released, we don't know whether the locked rows match our WHERE clause filter.

So we'll wait. Forever.



2.4 p10

Let's try another query.

You only have the clustered index on ID,
the white pages of the table.

What's the execution plan for this query:

```
SELECT *
FROM dbo.Users
WHERE Id = 26837
```

dbo.Users - Clustered Index

| Id | Rep | CreationDate | DisplayName | LastAccessDate | Location | Age | AboutMe | | |
|-----|------|------------------|--------------|------------------|----------------|------|-----------------------------|--|--|
| 1 | 2406 | 7/12/09 10:51 PM | Jeff Atwood | 4/1/10 10:35 AM | El Cerrito, CA | 39 | = '2013/11/10'
11 AND LastAccessDate <= '2013/11/11'
12 GO
13
14
```

The right window contains a SELECT statement:

```
1 SELECT *
2 FROM dbo.Users
3 WHERE Id = 26837
4 GO
5
6
7
8
9
10
11
12
13
14
```

Both windows have 'Messages' and 'Execution plan' tabs selected. The 'Execution plan' tab for the SELECT statement shows a 'Clustered Index Seek (Clustered)' operation with a cost of 100%.

Some of the rows of the clustered index (white pages) are locked, but not Brent's row.

Can Brent seek in and get unlocked data?



2.4 p12

## **Ways to work around the blocking**

- Add NOLOCK to our SELECT query
- Commit our UPDATE transaction faster
- Enable Read Committed Snapshot Isolation (RCSI)
- Add an index on LastAccessDate



2.4 p13

## Let's index LastAccessDate.

Nonclustered indexes are like separate copies of the table, with just the fields we want.

Cancel (roll back) the update, and create this index:

```
CREATE INDEX
IX_LastAccessDate_Id
ON dbo.Users(LastAccessDate, Id)
```

dbo.Users - IX\_LastAccessDate

| LastAccessDate   | Id | LastAccessDate   | Id  | LastAccessDate   | Id  | LastAccessDate   | Id   |
|------------------|----|------------------|-----|------------------|-----|------------------|------|
| 7/31/08 12:00 AM | -1 | 7/15/09 8:53 AM  | 445 | 7/15/09 9:10 PM  | 200 | 8/11/09 7:17 PM  | 39   |
| 7/15/09 7:08 AM  | 22 | 7/15/09 8:58 AM  | 457 | 7/16/09 6:22 AM  | 678 | 8/12/09 2:54 PM  | 943  |
| 7/15/09 7:10 AM  | 33 | 7/15/09 9:17 AM  | 501 | 7/17/09 2:30 AM  | 131 | 8/13/09 4:26 PM  | 364  |
| 7/15/09 7:11 AM  | 40 | 7/15/09 9:28 AM  | 524 | 7/17/09 9:30 AM  | 297 | 8/15/09 5:03 PM  | 910  |
| 7/15/09 7:11 AM  | 41 | 7/15/09 9:30 AM  | 527 | 7/17/09 8:43 PM  | 998 | 8/17/09 8:42 AM  | 202  |
| 7/15/09 7:11 AM  | 44 | 7/15/09 9:58 AM  | 587 | 7/18/09 12:38 PM | 394 | 8/17/09 10:11 AM | 628  |
| 7/15/09 7:12 AM  | 52 | 7/15/09 10:00 AM | 594 | 7/18/09 2:15 PM  | 924 | 8/17/09 10:33 AM | 157  |
| 7/15/09 7:13 AM  | 64 | 7/15/09 10:02 AM | 597 | 7/19/09 10:26 PM | 336 | 8/17/09 4:24 PM  | 1006 |

## Try your update – what's the plan?

```
BEGIN TRAN
UPDATE dbo.Users
SET Reputation = Reputation + 100
WHERE LastAccessDate >= '2013/11/10'
AND LastAccessDate <= '2013/11/11'
```

dbo.Users - Clustered Index

| Id  | Rep  | CreationDate     | DisplayName  | LastAccessDate   | Location       | Age  | AboutMe                      |  |  |
|-----|------|------------------|--------------|------------------|----------------|------|------------------------------|--|--|
| 1   | 2406 | 7/12/09 10:51 PM | Jeff Atwood  | 4/1/10 10:35 AM  | El Cerrito, CA | 39   | = '1800/01/01'  
AND LastAccessDate <= '1800/01/02'
```

dbo.Users - IX_LastAccessDate

| LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id |
|------------------|----|------------------|-----|------------------|-----|------------------|------|
| 7/31/08 12:00 AM | -1 | 7/15/09 8:53 AM | 445 | 7/15/09 9:10 PM | 200 | 8/11/09 7:17 PM | 39 |
| 7/15/09 7:08 AM | 22 | 7/15/09 8:58 AM | 457 | 7/16/09 6:22 AM | 678 | 8/12/09 2:54 PM | 943 |
| 7/15/09 7:10 AM | 33 | 7/15/09 9:17 AM | 501 | 7/17/09 2:30 AM | 131 | 8/13/09 4:26 PM | 364 |
| 7/15/09 7:11 AM | 40 | 7/15/09 9:28 AM | 524 | 7/17/09 9:30 AM | 297 | 8/15/09 5:03 PM | 910 |
| 7/15/09 7:11 AM | 41 | 7/15/09 9:30 AM | 527 | 7/17/09 8:43 PM | 998 | 8/17/09 8:42 AM | 202 |
| 7/15/09 7:11 AM | 44 | 7/15/09 9:58 AM | 587 | 7/18/09 12:38 PM | 394 | 8/17/09 10:11 AM | 628 |
| 7/15/09 7:12 AM | 52 | 7/15/09 10:00 AM | 594 | 7/18/09 2:15 PM | 924 | 8/17/09 10:33 AM | 157 |
| 7/15/09 7:13 AM | 64 | 7/15/09 10:02 AM | 597 | 7/19/09 10:26 PM | 336 | 8/17/09 4:24 PM | 1006 |

The SELECT finishes immediately.

Presto! The black pages weren't locked.

```
SQLQuery1.sql* 22
23 BEGIN TRAN
24 UPDATE dbo.Users
25     SET Reputation = Reputation + 100
26     WHERE LastAccessDate >= '2013/11/10'
27     AND LastAccessDate <= '2013/11/11';
28 GO
100 % 28
Messages Execution plan
(1330 row(s) affected)
(1 row(s) affected)

SQLQuery2.sql* 1
1 SELECT Id
2 FROM dbo.Users
3 WHERE LastAccessDate >= '1800/01/01'
4 AND LastAccessDate <= '1800/01/02'
5
100 %
Results Messages Execution plan
Query 1: Query cost (relative to the batch):
SELECT [Id] FROM [dbo].[Users] WHERE [LastAccessDate] >= '1800/01/01' AND [LastAccessDate] <= '1800/01/02'
Index Seek (NonClustered)
[Users].[IX_LastAccessDate_Id]
Cost: 100
SELECT Cost: 0
```

2.4 p19

The SELECT finishes immediately.

```
22
23  BEGIN TRAN
24  UPDATE dbo.Users
25      SET Reputation = Reputation + 100
26      WHERE LastAccessDate >= '2013/11/10'
27          AND LastAccessDate <= '2013/11/11';
28  GO
```

```
1  SELECT Id
2  FROM dbo.Users
3  WHERE LastAccessDate >= '1800/01/01'
4  AND LastAccessDate <= '1800/01/02'
5
```

100 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch):
SELECT [Id] FROM [dbo].[Users] WHERE [LastAccessDate] >= '1800/01/01' AND [LastAccessDate] <= '1800/01/02'

Index Seek (NonClustered)
[Users].[IX_LastAccessDate_Id]
Cost: 100

Presto! The black pages weren't locked.

But what if we query the same dates that we're updating?

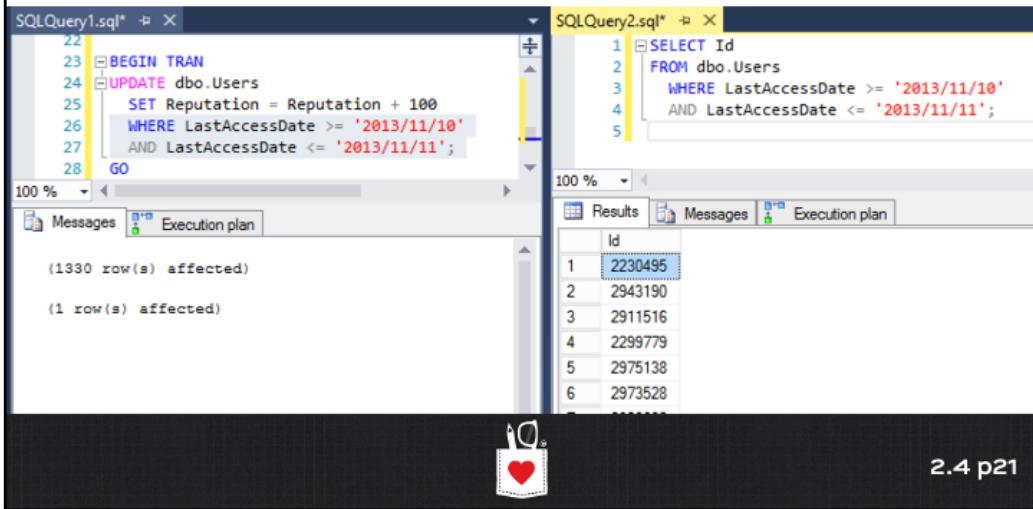


2.4 p2O

Still no blocking.

Our SELECT finishes instantly.

Indexes are amazing and the cure to all ills.



SQLQuery1.sql* 22
23 BEGIN TRAN
24 UPDATE dbo.Users
25 SET Reputation = Reputation + 100
26 WHERE LastAccessDate >= '2013/11/10'
27 AND LastAccessDate <= '2013/11/11';
28 GO

100 %

Messages Execution plan

(1330 row(s) affected)

(1 row(s) affected)

SQLQuery2.sql* 1
2 SELECT Id
3 FROM dbo.Users
4 WHERE LastAccessDate >= '2013/11/10'
5 AND LastAccessDate <= '2013/11/11';

100 %

Results Messages Execution plan

| Id |
|-----------|
| 1 2230495 |
| 2 2943190 |
| 3 2911516 |
| 4 2299779 |
| 5 2975138 |
| 6 2973528 |

2.4 p21

Now try this SELECT query.

What's the execution plan for this query:

```
SELECT Id, Reputation
FROM dbo.Users
WHERE LastAccessDate >= '1800/01/01'
AND LastAccessDate <= '1800/01/02'
```

dbo.Users - IX_LastAccessDate

| LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id |
|------------------|----|------------------|-----|------------------|-----|------------------|------|
| 7/31/08 12:00 AM | -1 | 7/15/09 8:53 AM | 445 | 7/15/09 9:10 PM | 200 | 8/11/09 7:17 PM | 39 |
| 7/15/09 7:08 AM | 22 | 7/15/09 8:58 AM | 457 | 7/16/09 6:22 AM | 678 | 8/12/09 2:54 PM | 943 |
| 7/15/09 7:10 AM | 33 | 7/15/09 9:17 AM | 501 | 7/17/09 2:30 AM | 131 | 8/13/09 4:26 PM | 364 |
| 7/15/09 7:11 AM | 40 | 7/15/09 9:28 AM | 524 | 7/17/09 9:30 AM | 297 | 8/15/09 5:03 PM | 910 |
| 7/15/09 7:11 AM | 41 | 7/15/09 9:30 AM | 527 | 7/17/09 8:43 PM | 998 | 8/17/09 8:42 AM | 202 |
| 7/15/09 7:11 AM | 44 | 7/15/09 9:58 AM | 587 | 7/18/09 12:38 PM | 394 | 8/17/09 10:11 AM | 628 |
| 7/15/09 7:12 AM | 52 | 7/15/09 10:00 AM | 594 | 7/18/09 2:15 PM | 924 | 8/17/09 10:33 AM | 157 |
| 7/15/09 7:13 AM | 64 | 7/15/09 10:02 AM | 597 | 7/19/09 10:26 PM | 336 | 8/17/09 4:24 PM | 1006 |

Your execution plan:

1. Use the new index on LastAccessDate – seek directly to 1800/01/01, make a list of rows that match. (There won't be any, right?)
2. Look up their IDs in the clustered index (white pages), and get their Reputation field

Will it ~~blend~~ be blocked?

dbo.Users - IX_LastAccessDate

| LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id | LastAccessDate | Id |
|------------------|----|------------------|-----|------------------|-----|------------------|------|
| 7/31/08 12:00 AM | -1 | 7/15/09 8:53 AM | 445 | 7/15/09 9:10 PM | 200 | 8/11/09 7:17 PM | 39 |
| 7/15/09 7:08 AM | 22 | 7/15/09 8:58 AM | 457 | 7/16/09 6:22 AM | 678 | 8/12/09 2:54 PM | 943 |
| 7/15/09 7:10 AM | 33 | 7/15/09 9:17 AM | 501 | 7/17/09 2:30 AM | 131 | 8/13/09 4:26 PM | 364 |
| 7/15/09 7:11 AM | 40 | 7/15/09 9:28 AM | 524 | 7/17/09 9:30 AM | 297 | 8/15/09 5:03 PM | 910 |
| 7/15/09 7:11 AM | 41 | 7/15/09 9:30 AM | 527 | 7/17/09 8:43 PM | 998 | 8/17/09 8:42 AM | 202 |
| 7/15/09 7:11 AM | 44 | 7/15/09 9:58 AM | 587 | 7/18/09 12:38 PM | 394 | 8/17/09 10:11 AM | 628 |
| 7/15/09 7:12 AM | 52 | 7/15/09 10:00 AM | 594 | 7/18/09 2:15 PM | 924 | 8/17/09 10:33 AM | 157 |
| 7/15/09 7:13 AM | 64 | 7/15/09 10:02 AM | 597 | 7/19/09 10:26 PM | 336 | 8/17/09 4:24 PM | 1006 |

It finishes instantly!

The screenshot shows two SQL queries in SSMS:

SQLQuery1.sql:

```
22
23 BEGIN TRAN
24 UPDATE dbo.Users
25     SET Reputation = Reputation + 100
26     WHERE LastAccessDate >= '2013/11/10'
27     AND LastAccessDate <= '2013/11/11';
28 GO
```

SQLQuery2.sql:

```
1 SELECT Id, Reputation
2     FROM dbo.Users
3     WHERE LastAccessDate >= '1800/01/01'
4     AND LastAccessDate <= '1800/01/01';
```

Execution plan for Query 2:

```
Query 1: Query cost (relative to the batch): 100%
SELECT [Id],[Reputation] FROM [dbo].[Users] WHERE [LastAccessDate] >= '1800/01/01' AND [LastAccessDate] <= '1800/01/01';

Nested Loops (Inner Join)
    Index Seek (NonClustered) [Users].[IX_LastAccessDate_Id]
        Cost: 50 %

    Key Lookup (Clustered) [Users].[PK_Users_Id]
        Cost: 50 %
```

Of course, no rows are returned.
But what happens if we look at 2013/11/10?



2.4 p24

Awww, shucks.

```
SQLQuery1.sql* - X
22
23 BEGIN TRAN
24 UPDATE dbo.Users
25     SET Reputation = Reputation + 100
26     WHERE LastAccessDate >= '2013/11/10'
27     AND LastAccessDate <= '2013/11/11';
28 GO
100 % < >
Messages Execution plan
(1330 row(s) affected)
(1 row(s) affected)

SQLQuery2.sql - Executing...* - X
1 SELECT Id, Reputation
2     FROM dbo.Users
3     WHERE LastAccessDate >= '2013/11/10'
4     AND LastAccessDate <= '2013/11/11';
5
100 % < >
Results Messages
```

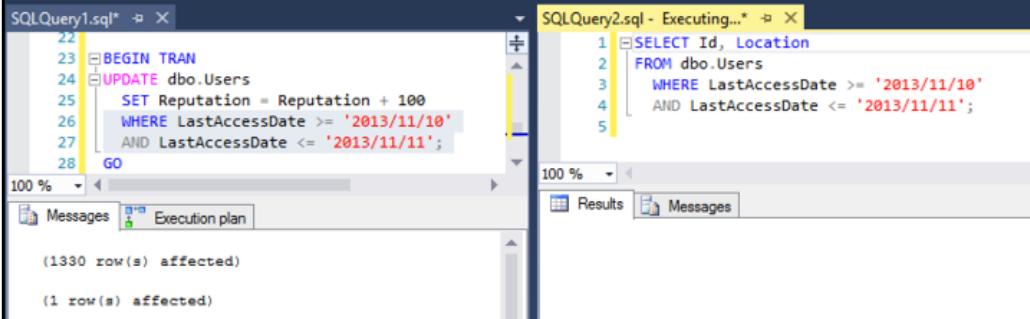
The **SELECT** on the right is blocked because we need Reputation, which is currently being edited.

But what if we skip Reputation, and get Location?



2.4 p25

It's still blocked.



```
SQLQuery1.sql* - Executing...  
22  
23 BEGIN TRAN  
24 UPDATE dbo.Users  
25     SET Reputation = Reputation + 100  
26     WHERE LastAccessDate >= '2013/11/10'  
27     AND LastAccessDate <= '2013/11/11';  
28 GO  
  
100 %  
Messages Execution plan  
  
(1330 row(s) affected)  
  
(1 row(s) affected)
```

```
SQLQuery2.sql - Executing...  
1 SELECT Id, Location  
2 FROM dbo.Users  
3 WHERE LastAccessDate >= '2013/11/10'  
4 AND LastAccessDate <= '2013/11/11';  
5
```

100 %

Results Messages

The lock is on the clustered index row, not specific fields.

So if I wanted to query Id and Location, while I was updating Reputation, and not get blocked, what could I do?



2.4 p26

Recap so far

SQL Server locks individual indexes at the row level at first, and only the relevant indexes – not all of ‘em.

Indexes are like readable replicas inside our DB.

Avoid indexing “hot” fields if you can.

Even just including “hot” fields comes at a price.



2.4 p27

Let's reward more people.

In your transaction window, let's add another update without committing it.

What's the execution plan for this query:

```
BEGIN TRAN  
UPDATE dbo.Users  
SET Reputation = Reputation + 100  
WHERE LastAccessDate >= '2014/11/10'  
AND LastAccessDate <= '2014/11/11'
```

dbo.Users - Clustered Index

| ID | Rep | CreationDate | DisplayName | LastAccessDate | Location | Age | AboutMe |
|-----|------|------------------|--------------|------------------|----------------|------|------------------------------|
| 1 | 2406 | 7/12/09 10:51 PM | Jeff Atwood | 4/1/10 10:35 AM | El Cerrito, CA | 39 | X
25 GO
26 UPDATE dbo.Users
27 SET Reputation = Reputation + 100
28 WHERE LastAccessDate > "2014/11/10"
29 AND LastAccessDate <= "2014/11/11";
30 GO

SQLQuery2.sql" -> X
1 SELECT *
2 FROM dbo.Users
3 WHERE Id = 26857;
4

Results Messages Execution plan
```

| ID | AcceptMe | Age | CreationDate            | DisplayName | DownVotes | EmailHash | LastAccessDate          |
|----|----------|-----|-------------------------|-------------|-----------|-----------|-------------------------|
| 1  | 26857    | 43  | 2008-10-10 14:26:33.540 | Brent Ozar  | 12        | NULL      | 2016-03-02 20:16:02.267 |

1 row(s) affected  
(1 row(s) affected)



2.4 p31

## Let's do one more round of gifts.

In your transaction window, let's add another update without committing it.

What's the execution plan for this query:

```
BEGIN TRAN
UPDATE dbo.Users
SET Reputation = Reputation + 100
WHERE LastAccessDate >= '2015/11/10'
AND LastAccessDate <= '2015/11/11'
```

dbo.Users - Clustered Index

| ID  | Rep  | CreationDate     | DisplayName  | LastAccessDate   | Location       | Age  | AboutMe                      |
|-----|------|------------------|--------------|------------------|----------------|------|------------------------------|
| 1   | 2406 | 7/12/09 10:51 PM | Jeff Atwood  | 4/1/10 10:35 AM  | El Cerrito, CA | 39   | = '2015/11/10'
37     AND LastAccessDate <= '2015/11/11';
38 GO
39

SQLQuery2.sql - Executing...*  # X
1 SELECT *
2     FROM dbo.Users
3     WHERE Id = 26837;
4

100 %  Results  Messages

Messages Execution plan

(6407 row(s) affected)
(1 row(s) affected)
```

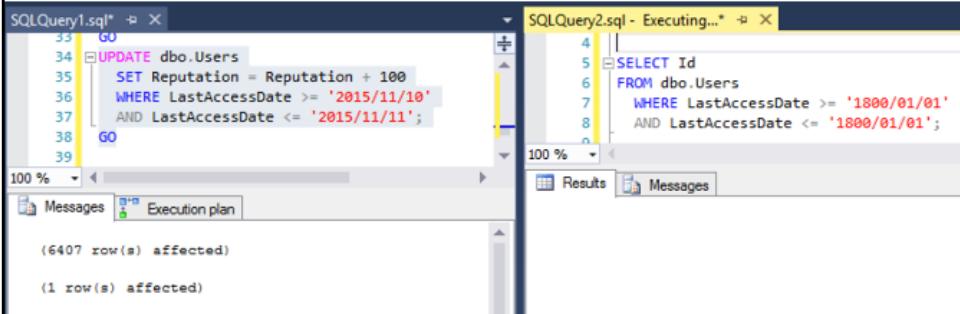
Brent shouldn't be blocked, but this SELECT hangs.

SQL Server has gone from locking individual rows of the clustered index, up to locking the entire index.



2.4 p34

Can we query by LastAccessDate?



```
SQLQuery1.sql - Executing... - X
33 GO
34 UPDATE dbo.Users
35     SET Reputation = Reputation + 100
36     WHERE LastAccessDate >= '2015/11/10'
37     AND LastAccessDate <= '2015/11/11';
38 GO
39

SQLQuery2.sql - Executing... - X
4 SELECT Id
5 FROM dbo.Users
6 WHERE LastAccessDate >= '1800/01/01'
7     AND LastAccessDate <= '1800/01/01';

100 %
```

Results Messages

(6407 row(s) affected)

(1 row(s) affected)

Nope, that's blocked too now, even for 1800.

Our row-level locks got escalated to table locks.



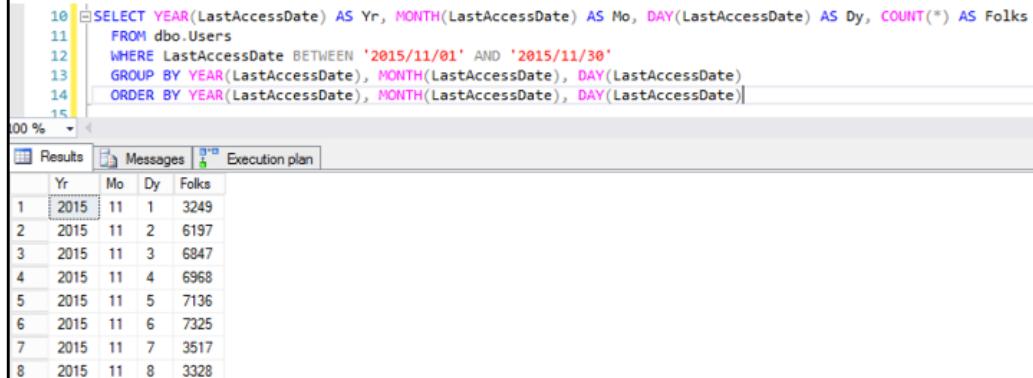
2.4 p35

That escalated quickly

SQL Server needs memory to track locks.

When queries hold thousands of row-level locks,
SQL Server escalates those locks to table-level.

Depending on what day(s) of data you're updating,
you might get row-level or table-level locks.



A screenshot of SQL Server Management Studio (SSMS) showing a query results grid. The query is a SELECT statement that groups data by year, month, and day, and counts the number of users (Folks) for each group. The results show data for November 2015, with the first row being highlighted.

| | Yr | Mo | Dy | Folks |
|---|------|----|----|-------|
| 1 | 2015 | 11 | 1 | 3249 |
| 2 | 2015 | 11 | 2 | 6197 |
| 3 | 2015 | 11 | 3 | 6847 |
| 4 | 2015 | 11 | 4 | 6968 |
| 5 | 2015 | 11 | 5 | 7136 |
| 6 | 2015 | 11 | 6 | 7325 |
| 7 | 2015 | 11 | 7 | 3517 |
| 8 | 2015 | 11 | 8 | 3328 |

What we learned so far

Indexes aren't just great for selects:
they can make DUI operations faster, too.

You want the right number of indexes to support all of
your concurrent operations, but no more than that.

Be aware that lock escalation stops queries dead.

Before tuning specific queries, use sp_BlitzIndex to:

- Remove indexes that aren't getting used
- Put a clustered index on OLTP tables w/DUIs
- Add high-value indexes that are really needed



2.4 p37

Tools to find & reduce blocking

The D.E.A.T.H. Method:

- D.E.A. with `sp_BlitzIndex`:
Look for “Aggressive Indexes” warnings
- T. with `sp_BlitzCache @SortOrder = ‘duration’`
Look for “Long Running, Low CPU” warnings
- H. because heaps can be locking hell

Finding deadlocks: `sp_BlitzLock`

Finding queries live: `sp_BlitzWho`, `sp_WhoIsActive`

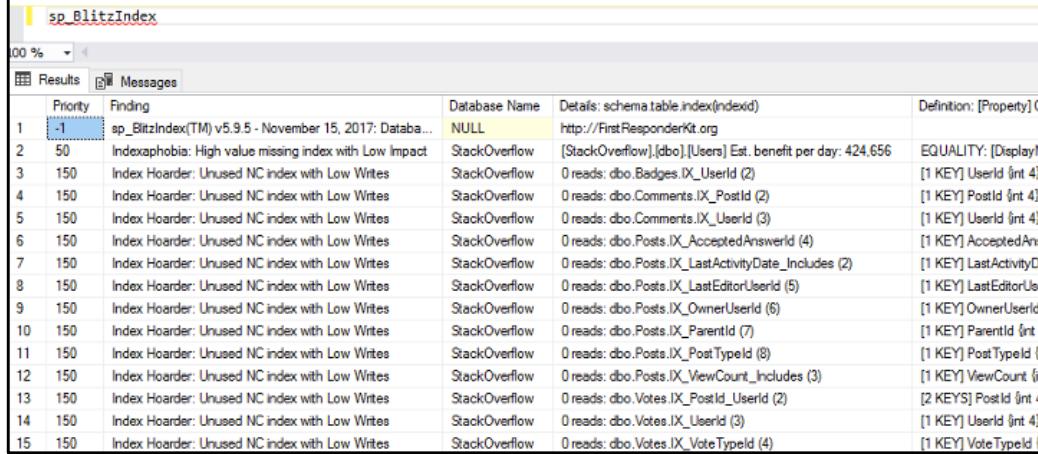


2.4 p38

About sp_BlitzIndex

Totally free index health check from BrentOzar.com

Gives your indexes a psychiatric exam



| | Priority | Finding | Database Name | Details: schema table index(indexid) | Definition: [Property] C |
|----|----------|---|---------------|---|-------------------------------------|
| 1 | -1 | sp_BlitzIndex(TM) v5.9.5 - November 15, 2017: Database... | NULL | http://FirstResponderKit.org | |
| 2 | 50 | Indexaphobia: High value missing index with Low Impact | StackOverflow | [StackOverflow].[dbo].[Users] | Est. benefit per day: 424.656 |
| 3 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Badges.IX_Userid (2) | [1 KEY] Userid (int 4) |
| 4 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Comments.IX_PostId (2) | [1 KEY] PostId (int 4) |
| 5 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Comments.IX_Userid (3) | [1 KEY] Userid (int 4) |
| 6 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_AcceptedAnswerId (4) | [1 KEY] AcceptedAnswerId (int 4) |
| 7 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_LastActivityDate_Includes (2) | [1 KEY] LastActivityDate (datetime) |
| 8 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_LastEditorUserId (5) | [1 KEY] LastEditorUserId (int 4) |
| 9 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_OwnerUserId (6) | [1 KEY] OwnerUserId (int 4) |
| 10 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_ParentId (7) | [1 KEY] ParentId (int 4) |
| 11 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_PostTypeId (8) | [1 KEY] PostTypeId (int 4) |
| 12 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Posts.IX_ViewCount_Includes (3) | [1 KEY] ViewCount (int 4) |
| 13 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Votes.IX_PostId_Userid (2) | [2 KEYS] PostId (int 4) |
| 14 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Votes.IX_Userid (3) | [1 KEY] Userid (int 4) |
| 15 | 150 | Index Hoarder: Unused NC index with Low Writes | StackOverflow | 0 reads: dbo.Votes.IX_VoteTypeId (4) | [1 KEY] VoteTypeId (int 4) |

“Aggressive Indexes” warning

Aggressive Indexes warning:

means SQL Server is tracking blocking on this index.

That doesn't mean this index is the problem one, it's just where the locking is happening.

Aggressive Indexes, Under-Indexed:

the table has 0-3 nonclustered indexes, and probably needs an index to support the D/U/I operations.

Aggressive Indexes, Over-Indexed:

the table has 10+ indexes, probably needs a haircut, especially indexes with writes > 0, but 0 reads.

They're only slowing you down.



2.4 p40

You probably won't see “Aggressive Indexes” in labs.

If you do labs to follow along with training, you probably won't get “Aggressive Indexes” warnings.

They're triggered off fairly high volumes of blocking, more than we'll experience in a few hours of queries.

If you *do* see them, you won't be able to make them go away: this DMV doesn't reset until your SQL Server instance is restarted or the index is dropped & recreated.



2.4 p41

We focused on #1.

1. Have enough indexes to make your queries fast, but not so many that they slow down DUIs, making them hold more locks for longer times.
(This session focuses on this.)
2. Keep batches & transactions short and sweet.
(We cover this in Mastering Query Tuning.)
3. Use the right isolation level for your app's needs.
(We cover this in Mastering Server Tuning.)



2.4 p42

#2: keeping batches & transactions short and sweet

Make a decision:

- Lock the whole table and get done fast, or
- Work in small batches, avoiding table locks

Then design your code to:

- Be sargable, so SQL Server can predict how many rows will be affected
- Avoid batch sizes set by variables (TOP @i)
- Move things out of the transaction where possible
- Work through tables in a predictable order



2.4 p43

#3: consider optimistic locking

SQL Server defaults to pessimistic locking

You have other options:

- Snapshot isolation
- Read committed snapshot isolation (RCSI)

Other platforms default to RCSI (Oracle, Azure SQL)

Readers don't block writers, and
writers don't block readers

It does have drawbacks though, beyond what I can
cover fast. Learn more: BrentOzar.com/go/rcsi



2.4 p44

Related learning

Take Care When Scripting Batches by Michael J. Swart: <https://michaeljswart.com/2014/09/take-care-when-scripting-batches/>

Which Locks Count Toward Lock Escalation by Kendra Little:
<https://littlekendra.com/2017/04/03/which-locks-count-toward-lock-escalation/>

Video: Using Batches to Do A Lot of Work Without Blocking by me (in Mastering Query Tuning class)



2.4 p45